

# Dongsheng Ding

---

CONTACT INFORMATION      Department of Electrical and Systems Engineering      Phone: (213) 574-9471  
University of Pennsylvania      E-mail: [dongshed@seas.upenn.edu](mailto:dongshed@seas.upenn.edu)  
3333 Chestnut Street, Philadelphia, PA 19104      URL: <https://dongshed.github.io>

TEACHING INTERESTS      I view teaching as a responsibility toward future generations and society. Technology rapidly advances, yet engineering course materials stay relatively static in traditional fields, limiting alignment with today's interdisciplinary applications and innovations. My teaching aims to develop forward-thinking curricula that bridge interdisciplinary gaps and expand career possibilities.

RESEARCH INTERESTS      Application of principles and tools in optimization and control to the *study of requirement-driven machine decision-making*.

Topics of interest span: constrained reinforcement learning; multi-agent reinforcement learning; constrained generative modeling; constrained alignment of diffusion and large language models; robustness and resilience analysis of optimization algorithms; nonlinear control of fractional-order dynamical systems.

Applications of interest include: autonomous robots/vehicles; generative artificial intelligence in applied science and engineering, such as natural language processing, image processing, and molecular design; decision-making systems for power grid, transportation, communication, and healthcare management.

APPOINTMENT      **Postdoctoral Researcher**      September 2022 – Now  
University of Pennsylvania

**Graduate Research Assistant**      August 2017 – August 2022  
University of Southern California

**Graduate Research Assistant & Teaching Assistant**      August 2015 – August 2017  
University of Minnesota, Twin Cities

EDUCATION      **PhD in Electrical Engineering**, GPA: 4.0/4.0      August 2022  
University of Southern California  
Thesis: Provable reinforcement learning for constrained and multi-agent control systems  
Advisor: Mihailo R. Jovanović

**MS in Electrical Engineering**, GPA: 3.9/4.0      August 2017  
University of Minnesota, Twin Cities

**ME in Control Theory & Engineering**, GPA: 3.7/4.0      March 2015  
**BE in Automation**, GPA: 3.8/4.0      June 2011  
Zhejiang University, Hangzhou, China

SUBMITTED  
PAPERS

1. B. Zhang, S. Li, I. Hounie, O. Bastani, D. Ding, A. Ribeiro. “Alignment of large language models with constrained learning,” 2025. (under review)
2. S. Khalafi, I. Hounie, D. Ding, A. Ribeiro. “Composition and alignment of diffusion models using constrained learning,” 2025. (under review)

REFEREED  
PUBLICATIONS

**Journals**

1. D. Ding, K. Zhang, J. Duan, T. Başar, and M. R. Jovanović. “Convergence and sample complexity of natural policy gradient primal-dual methods for constrained MDPs,” *J. Mach. Learn. Res.*; also arXiv: 2206.02346, 2022. (under review)
2. D. Ding, X. Wei, Z. Yang, Z. Wang, and M. R. Jovanović. “Fast multi-agent temporal-difference learning via homotopy stochastic primal-dual optimization,” *IEEE Trans. Control. Netw. Syst.*; also arXiv:1908.02805, 2020. (under revision)
3. Q. Wang, J. Zhang, D. Ding, and D. Qi, “Adaptive Mittag-Leffler stabilization of a class of fractional order uncertain nonlinear systems,” *Asian J. Control*, 18(6) 2343–2351, 2016.
4. D. Ding, D. Qi, and Q. Wang, “Asymptotic pseudo-state stabilization of uncertain fractional-order nonlinear systems with additive disturbance,” *Nonlinear Dyn.*, 81(1) 667–677, 2015.
5. Q. Wang, D. Ding, and D. Qi, “Mittag-Leffler synchronization of uncertain fractional order chaotic systems,” *Chinese Physics B*, 24(6), 2015.
6. D. Ding, D. Qi, and Q. Wang, “Nonlinear Mittag-Leffler stabilization of commensurate fractional-order nonlinear systems,” *IET Control Theory Appl.*, 9(5) 681–690, 2014.
7. D. Ding, D. Qi, X. Luo, J. Chen, X. Wang, and P. Du, “Convergence analysis and performance of an extended central force optimization algorithm,” *Appl. Math. Comput.*, 219(4), 2246–2259, 2012.
8. D. Ding, X. Luo, J. Chen, X. Wang, P. Du, and Y. Guo, “A convergence proof and parameter analysis of central force optimization algorithm,” *J. Convergence Inf. Technol.*, 6(10), 16–23, 2011.

**Machine Learning and Artificial Intelligence Conferences**

1. S. Rozada, D. Ding, A. Marques, A. Ribeiro. “Deterministic policy gradient primal-dual methods for continuous-space constrained MDPs.” in *Proceedings of the AAAI Conference on Artificial Intelligence*, Philadelphia, Pennsylvania, 2025; also arXiv:2408.10015. (**acceptance rate 23.4%**)
2. X. Huang, S. Li, E. Dobriban, O. Bastani, H. Hassani, and D. Ding. “One-shot safety alignment for large language models via optimal dualization.” in *Proceedings of the Advances in Neural Information Processing Systems*, Vancouver, Canada, 2024; also arXiv: 2405.19544. (**acceptance rate 25.8%, spotlight**)

3. S. Khalafi, D. Ding, and A. Ribeiro. “Constrained diffusion models via dual training.” in *Proceedings of the Advances in Neural Information Processing Systems*, Vancouver, Canada, 2024; also arXiv:2408.15094. (**acceptance rate 25.8%**)
4. D. Ding, Z. Huan, and A. Ribeiro. “Resilient constrained reinforcement learning.” in *Proceedings of the 27th International Conference on Artificial Intelligence and Statistics*, Valencia, Spain, 2024; also arXiv:2312.17194. (**acceptance rate 27.6%**)
5. D. Ding, C.-Y. Wei, K. Zhang, and A. Ribeiro. “Last-iterate convergent policy gradient primal-dual methods for constrained MDPs,” in *Proceedings of the Advances in Neural Information Processing Systems*, New Orleans, Louisiana, 2023; also arXiv:2306.11700. (**acceptance rate 26.1%**)
6. D. Ding, X. Wei, Z. Yang, Z. Wang, and M. R. Jovanović. “Sample efficient generalized Lagrangian policy optimization for safe multi-agent reinforcement learning,” in *Proceedings of the Learning for Dynamics and Control Conference*, Philadelphia, Pennsylvania, 2023; also arXiv:2306.00212.
7. D. Ding, C.-Y. Wei, K. Zhang, and M. R. Jovanović. “Independent policy gradient for large-scale Markov potential games: sharper rates, function approximation, and game-agnostic convergence,” in *Proceedings of the 39th International Conference on Machine Learning*, Baltimore, Maryland, 2022; also arXiv:2202.04129. (**acceptance rate 21.5%, 118/1117 long presentations**)
8. D. Ding, X. Wei, Z. Yang, Z. Wang, and M. R. Jovanović. “Provably efficient safe exploration via primal-dual policy optimization,” in *Proceedings of the 24th International Conference on Artificial Intelligence and Statistics*, Virtual, 2021; also arXiv:2003.00534. (**acceptance rate 30%, 48/455 orals**)
9. D. Ding, K. Zhang, T. Başar, and M. R. Jovanović. “Natural policy gradient primal-dual method for constrained Markov decision processes,” in *Proceedings of the Advances in Neural Information Processing Systems*, Virtual, 2020. (**acceptance rate 20%**)
10. D. Ding, X. Wei, Z. Yang, Z. Wang, and M. R. Jovanović. “Fast multi-agent temporal-difference learning via homotopy stochastic primal-dual method,” in *the Optimization Foundations for Reinforcement Learning Workshop at NeurIPS*, Vancouver, Canada, 2019; also arXiv:1908.02805.

### Control Conferences

1. D. Ding and M. R. Jovanović. “Policy gradient primal-dual mirror descent for constrained MDPs with large state spaces,” in *Proceedings of the 61st IEEE Conference on Decision and Control*, Cancún, Mexico, 2022.
2. D. Ding, K. Zhang, T. Basar and M. R. Jovanović, “Convergence and optimality of policy gradient primal-dual method for constrained Markov decision processes,” in *Proceedings of the 2022 American Control Conference*, Atlanta, Georgia, 2022.

3. D. Ding, X. Wei, H. Yu, and M. R. Jovanović. “Byzantine-resilient distributed learning under constraints,” in *Proceedings of the 2021 American Control Conference*, Virtual, 2021.
4. D. Ding, J. Yuan, and M. R. Jovanović. “Discounted online Newton method for time-varying time series prediction,” in *Proceedings of the 2021 American Control Conference*, Virtual, 2021.
5. D. Ding and M. R. Jovanović. “Global exponential stability of primal-dual gradient flow dynamics based on the proximal augmented Lagrangian: A Lyapunov-based approach,” in *Proceedings of the 59th IEEE Conference on Decision and Control*, Virtual, 2020.
6. D. Ding, X. Wei, and M. R. Jovanović. “Distributed robust statistical learning: Byzantine mirror descent,” in *Proceedings of the 58th IEEE Conference on Decision and Control*, Nice, France, 2019.
7. D. Ding and M. R. Jovanović. “Global exponential stability of primal-dual gradient flow dynamics based on the proximal augmented Lagrangian,” in *Proceedings of the 2019 American Control Conference*, Philadelphia, Pennsylvania, 2019.
8. D. Ding, B. Hu, N. K. Dhingra, and M. R. Jovanović. “An exponentially convergent primal-dual algorithm for nonsmooth composite minimization,” in *Proceedings of the 57th IEEE Conference on Decision and Control*, Miami Beach, Florida, 2018.
9. D. Ding and M. R. Jovanović. “A primal-dual Laplacian gradient flow dynamics for distributed resource allocation problems,” in *Proceedings of the 2018 American Control Conference*, Milwaukee, Wisconsin, 2018.
10. D. Ding, D. Qi, and Q. Wang, “Adaptive Mittag-Leffler stabilization of commensurate fractional-order nonlinear systems,” in *Proceedings of the 53rd IEEE Conference on Decision and Control*, Los Angeles, California, 2014.
11. D. Ding, G. Zhang, D. Qi, and H. Zhang, “Strategy analysis of an evolutionary spectrum sensing game,” in *Proceedings of the Intelligent Computing and Applications (LSMS & ICSEE)*, Shanghai, China, 2014. (**nominate paper award**)
12. D. Ding, D. Qi, and Q. Wang, “Alternative LMI characterizations for fractional-order linear systems,” in *Proceedings of the 33rd Chinese Control Conference*, Nanjing, China, 2014.
13. D. Ding, D. Qi, and Q. Wang, “Fractional-order integral state space modeling and quasi state analysis via block operational matrix scheme,” in *Proceedings of the 26th Chinese Control and Decision Conference*, Changsha, China, 2014.

## RESEARCH

## Invited Talks & Posters

## PRESENTATIONS

1. Invited talk of “Requirement-driven machine decision-making: From reinforcement learning to generative artificial intelligence,” in *the EECS Faculty Candidate Seminar*, University of Tennessee, Knoxville, 2025.

2. Invited talk of “Multi-agent reinforcement learning for large-scale Markov potential games,” in *the Workshop on Multi-Agent Learning in Dynamic Environments*, Texas A&M Institute of Data Science, TAMU, 2024.
3. Invited talk of “Multi-agent reinforcement learning for large-scale Markov potential games,” in *the ESE PhD Colloquium*, UPenn, 2024.
4. Invited talk of “Constrained policy optimization: A tale of regularization and optimism,” in *the INFORMS Optimization Society Conference*, Houston, Texas, 2024.
5. Invited talk of “Constrained policy optimization: A tale of regularization and optimism,” in *the ESE PhD Colloquium*, UPenn, 2023.
6. Invited talk of “Provable constrained policy optimization in reinforcement learning,” in *the Safe Reinforcement Learning Online Seminar*, Virtual, 2023.
7. Invited talk of “Finite-time performance of policy optimization methods for constrained reinforcement learning,” in *the INFORMS 2022 Annual Meeting*, Indianapolis, Indiana, 2022.
8. Invited poster of “Independent policy gradient for large-scale Markov potential games” in *the REAL@USC-Meta center workshop*, ECE, USC, 2022.
9. Invited poster of “Provably efficient safe exploration via primal-dual policy optimization” in *the 11th Annual Research Festival*, ECE, USC, 2021.
10. Invited talk of “Provable constrained policy optimization for reinforcement learning” in *the 38th Southern California Control Workshop*, University of California, Irvine, California, Virtual, 2021.
11. Invited poster of “Distributed robust statistical learning: Byzantine mirror descent” in *the 10th Annual Research Festival*, ECE, USC, 2019.
12. Invited poster of “An exponentially stable primal-dual algorithm for nonsmooth optimization” in *the 9th Annual Research Festival*, ECE, USC, 2018.
13. Invited talk of “A primal-dual algorithm for distributed resource allocation” in *the 34th Southern California Control Workshop*, University of California, Riverside, California, 2018.

### Conference Talks & Posters

1. Contributed poster of “One-shot safety alignment for large language models via optimal dualization.” in *the Next Generation of AI Safety Workshop at ICML and the Workshop on Theoretical Foundations of Foundation Models at ICML*, Vienna, Austria, 2024.
2. Contributed talk & poster of “Last-iterate convergent policy gradient primal-dual methods for constrained MDPs” in *the 37th Conference on Neural Information Processing Systems*, New Orleans, Louisiana, 2023.

3. Contributed poster of “Sample efficient generalized Lagrangian policy optimization for safe MARL” in *the 5th Annual Learning for Dynamics and Control Conference*, Philadelphia, Pennsylvania, 2023.
4. Contributed talk of “Policy gradient primal-dual mirror descent for constrained MDPs with large state spaces,” in *the 61st IEEE Conference on Decision and Control*, Cancún, Mexico, 2022.
5. Contributed talk of “Policy gradient primal-dual method for constrained MDPs,” in *the 2022 American Control Conference*, Atlanta, Georgia, 2022.
6. Contributed talk & poster of “Independent policy gradient for large-scale Markov potential games: sharper rates, function approximation, and game-agnostic convergence” in *the 39th International Conference on Machine Learning*, Baltimore, Maryland, 2022. (118/1117 long presentations)
7. Contributed talk & poster of “Provably efficient safe exploration via primal-dual policy optimization” in *the 24th International Conference on Artificial Intelligence and Statistics*, Virtual, 2021. (48/455 orals)
8. Contributed talk & poster of “Natural Policy Gradient Primal-Dual Method for Constrained Markov Decision Processes” in *the 34th Conference on Neural Information Processing Systems*, Virtual, 2020.
9. Contributed talk of “Global exponential stability of primal-dual gradient flow dynamics based on the proximal augmented Lagrangian” in *the 59th IEEE Conference on Decision and Control*, Virtual, 2020.
10. Contributed poster of “Fast multi-agent temporal-difference learning via homotopy stochastic primal-dual method,” in *the Optimization Foundations for Reinforcement Learning Workshop at NeurIPS*, Vancouver, Canada, 2019; *the Southern California Machine Learning Symposium*, UCSD, 2020.
11. Contributed talk of “Exponential stability of primal-dual gradient flow dynamics based on proximal augmented Lagrangian,” in *the 2019 American Control Conference*, Philadelphia, Pennsylvania, 2019.
12. Contributed talk of “Nonsmooth composite minimization: an exponentially convergent primal-dual algorithm,” in *the 57th IEEE Conference on Decision and Control*, Miami Beach, Florida, 2018.
13. Contributed talk of “A primal-dual Laplacian gradient flow dynamics for distributed resource allocation problems,” in *the 2018 American Control Conference*, Milwaukee, Wisconsin, 2018.
14. Contributed talk of “Adaptive Mittag-Leffler stabilization of commensurate fractional-order nonlinear systems” in *the 53rd IEEE Conference on Decision and Control*, Los Angeles, California, 2014.
15. Contributed talk of “Alternative LMI characterizations for fractional-order linear systems” in *the 33rd Chinese Control Conference*, Nanjing, China, 2014.

16. Contributed talk of “Fractional-order integral state space modeling” in *the 26th Chinese Control and Decision Conference*, Changsha, China, 2014.

HONORS & AWARDS	<b>Scholar Award</b> , Conference on Neural Information Processing Systems	2023
	<b>Expert Reviewers</b> , International Conference on Machine Learning	2021
	<b>Travel Award</b> , Conference on Neural Information Processing Systems	2020
	<b>Top Reviewers</b> , International Conference on Machine Learning	2020
	<b>Travel Award</b> , IEEE Conference on Decision and Control	2020
	<b>Travel Award</b> , American Control Conference	2018, 2019, 2022
	<b>MHI PhD Scholar Finalist</b> , ECE, University of Southern California	2018, 2021
	<b>ECE Department Fellowship</b> , University of Minnesota	2015
	<b>Honor for Outstanding Graduate Student</b> , Zhejiang University	2015
	<b>Nominate Paper Award</b> , LSMS & ICSEE, 2014, Shanghai	2014
	<b>Bosch Scholarship</b> , Bosch in China	2013
	<b>The First-Class of Graduate Scholarship</b> , Zhejiang University	2012–2015
	<b>National Scholarship</b> , Ministry of Education of P.R. China	2011
	<b>Wei Shaoxiang Engineering Talent</b> , Wei Shaoxiang Foundation, HongKong	2010
	<b>The Second-Class of Physics and Technology Innovation Contest</b> , Zhejiang Physical Society, Zhejiang, China	2009
	<b>The First-Class of Advanced Mathematics Contest</b> , Zhejiang Mathematical Society, Zhejiang, China	2008

ACADEMIC SERVICE	<b>Journal Referee</b>	
	IEEE Transactions on Control of Network Systems	
	IEEE Transactions on Automatic Control	
	IEEE Robotics and Automation Letters	
	IEEE Transactions on Intelligent Vehicles	
	IEEE Control Systems Letters	
	Systems & Control Letters	
	Optimization Letters	
	Machine Learning	
	Neural Networks	
	Automatica	
	IEEE Transactions on Pattern Analysis and Machine Intelligence	
	APSIPA Transactions on Signal and Information Processing	
	Engineering Applications of Artificial Intelligence	
	SIAM Journal on Mathematics of Data Science	
	Transactions on Machine Learning Research	
	Journal of Machine Learning Research	
	IET Control Theory & Applications	
	International Journal of Robust and Nonlinear Control	
	Frontiers of Information Technology & Electronic Engineering	
	International Journal of Systems Science	

The Journal of the Franklin Institute  
 Journal of Applied Mathematics and Computing  
 Nonlinear Dynamics  
 IEEE Access

#### Conference Referee

IEEE Conference on Decision and Control	2018–2024
Conference on Neural Information Processing Systems	2020–2025
International Conference on Learning Representations	2021–2025
International Conference on Artificial Intelligence and Statistics	2021, 2025
International Conference on Machine Learning	2020–2025
ICML Workshop on Next Generation of AI Safety	2024
ICML Workshop on Theoretical Foundations of Foundation Models	2024
NeurIPS Safe Generative AI Workshop	2025
Learning for Dynamics and Control Conference	2024, 2025
AAAI Conference on Artificial Intelligence	2023, 2024
American Control Conference	2018–2025
IFAC World Congress	2020
Chinese Control Conference	2014
Chinese Control and Decision Conference	2014

#### Conference Volunteer

38th International Conference on Machine Learning, Virtual	2021, 2022
24th International Conference on Artificial Intelligence and Statistics, Virtual	2021
35th Conference on Neural Information Processing Systems, Virtual	2021

#### Co-chair of Nonlinear System and Control Section

26th Chinese Control and Decision Conference, Changsha, China	2014
---	------

#### Admissions Committee, University of Pennsylvania

PhD Admission in the Department of Electrical and Systems Engineering	Fall 2023
---	-----------

#### TEACHING EXPERIENCE

<b>Lab Assistant</b> , University of Pennsylvania	
ESE 2000 <i>Artificial Intelligence Lab: Data, Systems, and Decisions</i>	Summer 2024
<b>Guest Lecturer</b> , University of Pennsylvania	
ESE 5140 <i>Graph Neural Networks</i>	Fall 2023
<b>Teaching Assistant</b> , University of Southern California	
EE 587 <i>Nonlinear Systems</i>	Spring 2018
<b>Guest Lecturer</b> , University of Minnesota Twin Cities	
EE 3015 <i>Statistical Methods in Electrical and Computer Engineering</i>	Spring 2017
<b>Teaching Assistant</b> , University of Minnesota Twin Cities	
EE 4231 <i>Automatic Control Systems</i>	Fall 2016
EE 3015 <i>Statistical Methods in Electrical and Computer Engineering</i>	Spring 2017
EE 8231 <i>Optimization Theory</i>	Spring 2017



MENTORING EXPERIENCE	<b>Mentor for Doctoral Students</b> , University of Pennsylvania	
	Shervin Khalafi (2nd year)	September 2023 – Now
	Topic: <i>Constrained learning for diffusion models</i>	
	Ignacio Hounie (4th year)	September 2024 – Now
	Topic: <i>Constrained learning for large language models</i>	
	Xinmeng Huang	February 2024 – February 2025
	PhD Thesis: <i>Topics in statistical machine learning: multitask learning, uncertainty quantification, and language model alignment</i>	
	Now: <i>Quantitative Researcher, Citadel Securities</i>	
	<b>Mentor for Doctoral Students</b> , King Juan Carlos University, Madrid, Spain	
	Sergio Rozada	February 2024 – February 2025
	PhD Thesis: <i>Low-rank methods in reinforcement learning</i>	
	<b>Mentor for Master Students</b> , University of Pennsylvania	
	Zhengyan Huan	September 2023 – February 2024
	Topic: Resilient constrained reinforcement learning	
	Now: <i>Doctoral Student, Tufts University</i>	
	<b>Mentor for Undergraduate Students</b> , University of Southern California	
	Terry Lu	January 2022 – June 2022
	Now: <i>Master in Finance, Princeton University</i>	
	<b>Academic Mentor for Master Students</b> , University of Southern California	
	Viterbi Graduate Mentorship Program	August 2018 – June 2021
SKILLS	Python, L <sup>A</sup> T <sub>E</sub> X, C/C++, Matlab, SQL	
MEMBERSHIPS	The Institute of Electrical and Electronics Engineers (IEEE)	
	IEEE Control Systems Society Membership	
	The Institute for Operations Research and the Management Sciences (INFORMS)	
	Stanford Encyclopedia of Philosophy	
GRADUATE COURSE HIGHLIGHTS		
	<b>Control Systems:</b> EE 5231 <i>Linear System and Optimal Control</i> , EE 8215 <i>Nonlinear Systems</i> , AEM 8421 <i>Robust Multi-Variable Control Systems</i> , AEM 8423 <i>Convex Optimization Methods in Control</i> ; <b>Optimization and Computation:</b> EE 5239 <i>Introduction to Nonlinear Optimization</i> , EE 8231 <i>Optimization Theory</i> , ISE 633 <i>Large-Scale Optimization for Machine Learning</i> , CSCI 5304 <i>Computational Aspects of Matrix Theory</i> , CSCI 8314 <i>Sparse Matrix Computations</i> ; <b>Probability and Machine Learning:</b> MATH 507A/B <i>Theory of Probability</i> , EE 556 <i>Stochastic Systems &amp; Reinforcement Learning</i> , CSCI 5525 <i>Machine Learning</i> , EE 546 <i>Mathematics of High-Dimensional Data</i> , DSO 699 <i>Statistical Learning Theory</i> , CSCI 699 <i>Theoretical Machine Learning</i> , EE 5581 <i>Information Theory and Coding</i> .	